

# Supervised Learning

John Joseph Valletta

University of Exeter, Penryn Campus, UK

April 2019



Researcher  
Development



# Overview

- What is supervised learning?
- Cross-validation
- $k$ -nearest neighbour ( $kNN$ )
- Decision trees
- Random forests
- Support vector machines

# What's the problem?

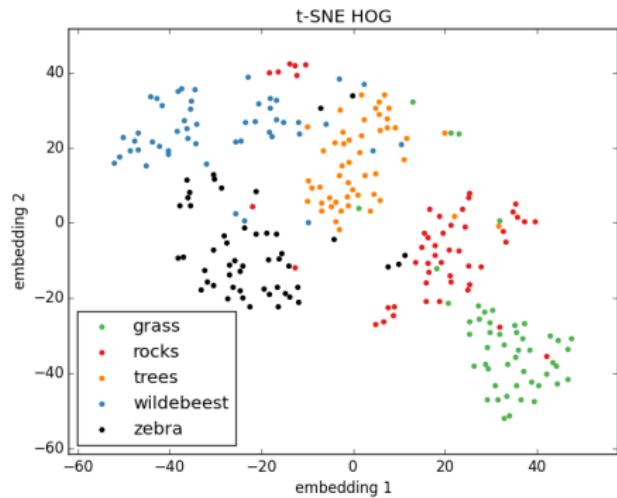


# What's the problem?



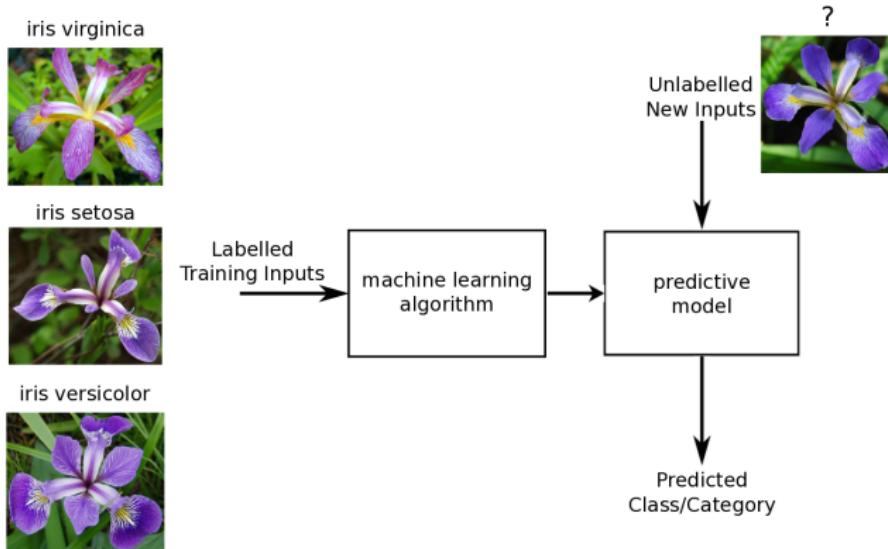
<https://www.livescience.com/23310-serengeti.html>

# Counting wildebeest



# What is supervised learning?

Supervised learning methods determine the mapping (**predictive model**) between a set of features and a continuous outcome (**regression**), or a categorical variable (**classification**)

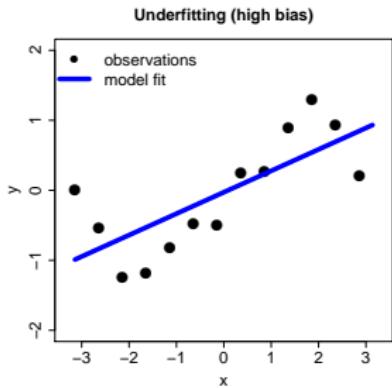


# Bias-variance tradeoff

- Machine learning algorithms are very flexible in order to deal with complex data
- So how *well* should we fit to the training data to get good generalisation?
- Driving the training error to zero is *not* a good idea

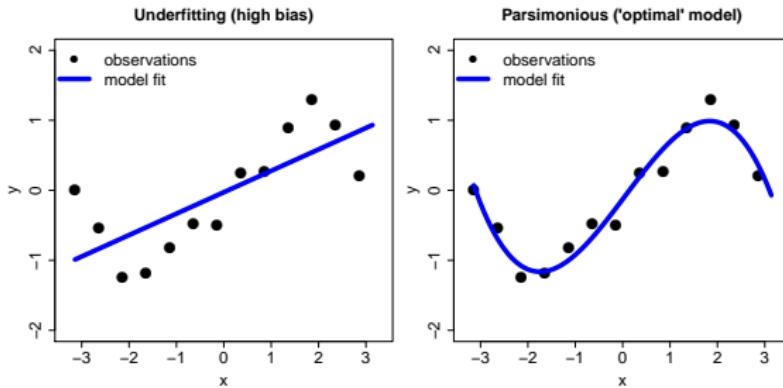
# Bias-variance tradeoff

- Machine learning algorithms are very flexible in order to deal with complex data
- So how *well* should we fit to the training data to get good generalisation?
- Driving the training error to zero is *not* a good idea



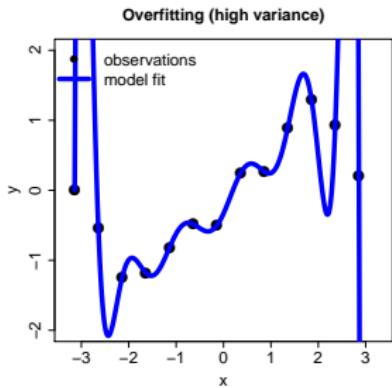
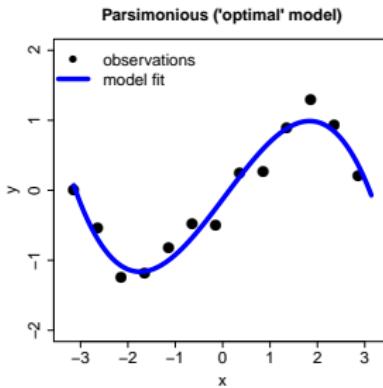
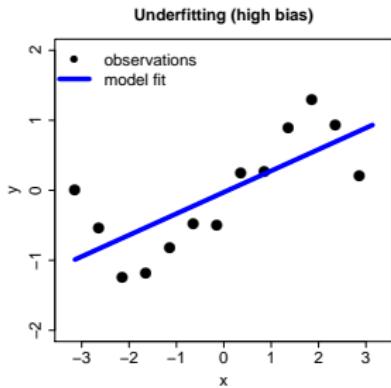
# Bias-variance tradeoff

- Machine learning algorithms are very flexible in order to deal with complex data
- So how *well* should we fit to the training data to get good generalisation?
- Driving the training error to zero is *not* a good idea



# Bias-variance tradeoff

- Machine learning algorithms are very flexible in order to deal with complex data
- So how *well* should we fit to the training data to get good generalisation?
- Driving the training error to zero is *not* a good idea

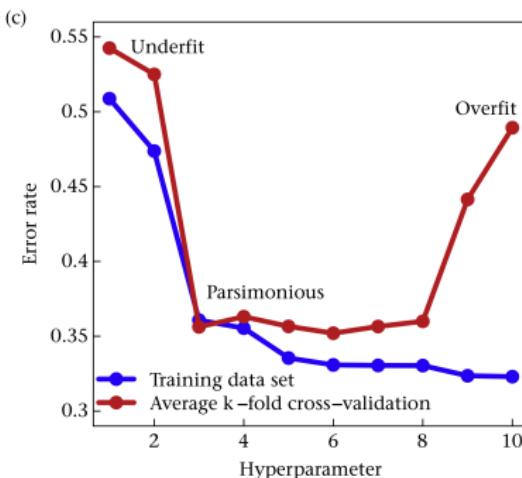
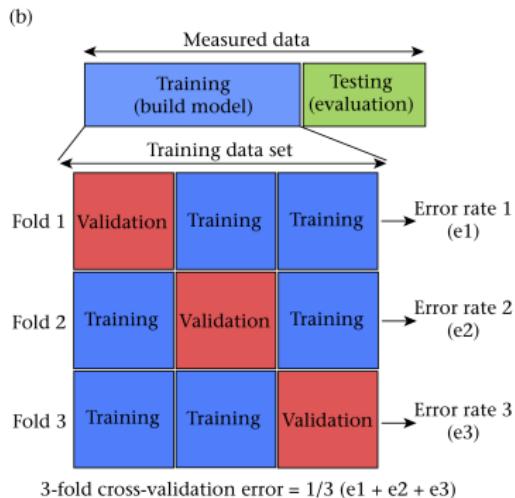


# Cross-validation

- The flexibility of ML models is constrained by **tuning** their **hyperparameters**
- ***k*-fold cross-validation** allow us to find optimal hyperparameters

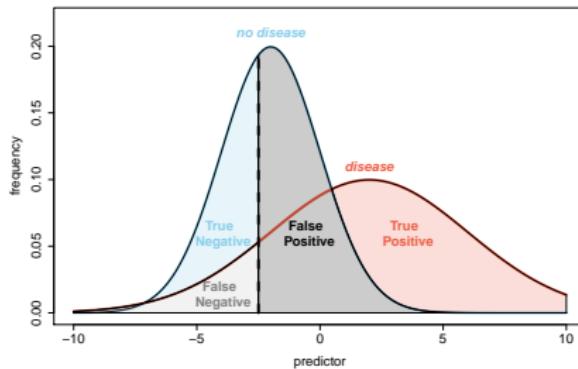
# Cross-validation

- The flexibility of ML models is constrained by **tuning** their **hyperparameters**
- k*-fold cross-validation** allow us to find optimal hyperparameters



# Predictive performance measures

- To compare models during cross-validation **predictive performance measures** are computed
- Several metrics exist, some of the more popular ones are:
  - **Regression:** root mean squared error (RMSE), R-squared
  - **Classification:** area under the ROC curve, confusion matrix



		Actual label				
		Grass	Rocks	Trees	Wildebeest	Zebra
Actual label	Grass	13	2	0	0	0
	Rocks	1	14	0	0	0
Trees	1	1	12	1	0	
Wildebeest	0	0	2	13	0	
Zebra	0	0	0	0	15	
		Predicted label				

# $k$ -nearest neighbour ( $k$ NN)

- ① Calculate distance between test point and every training data point
- ② Find the  $k$  training points closest to test point
- ③ Assign test point the majority vote of their class label

# $k$ -nearest neighbour ( $k$ NN)

- ① Calculate distance between test point and every training data point
- ② Find the  $k$  training points closest to test point
- ③ Assign test point the majority vote of their class label

# $k$ -nearest neighbour ( $k$ NN)

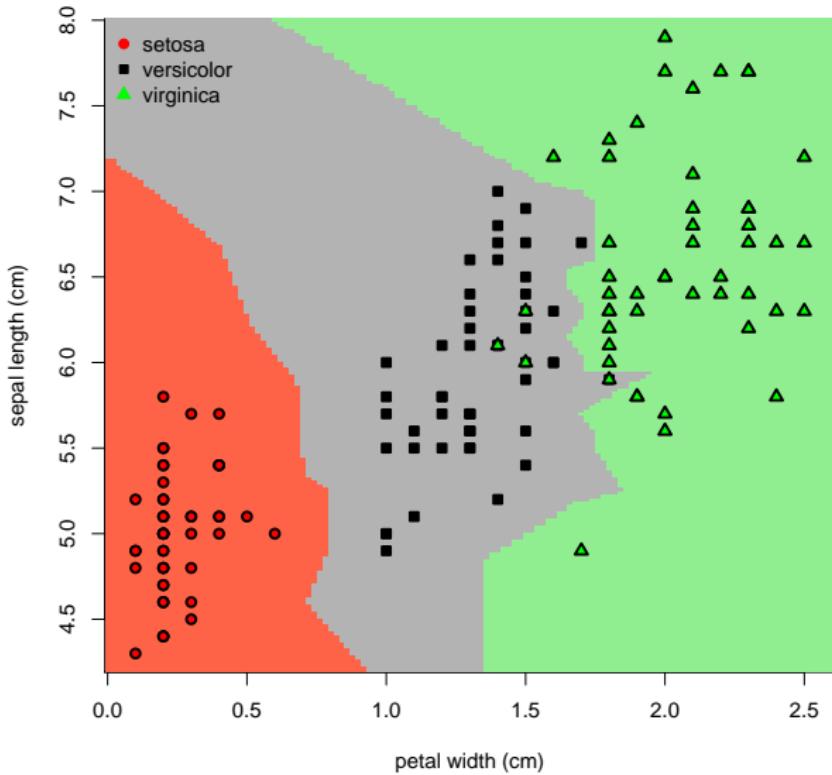
- ① Calculate distance between test point and every training data point
- ② Find the  $k$  training points closest to test point
- ③ Assign test point the majority vote of their class label

# $k$ -nearest neighbour ( $k$ NN)

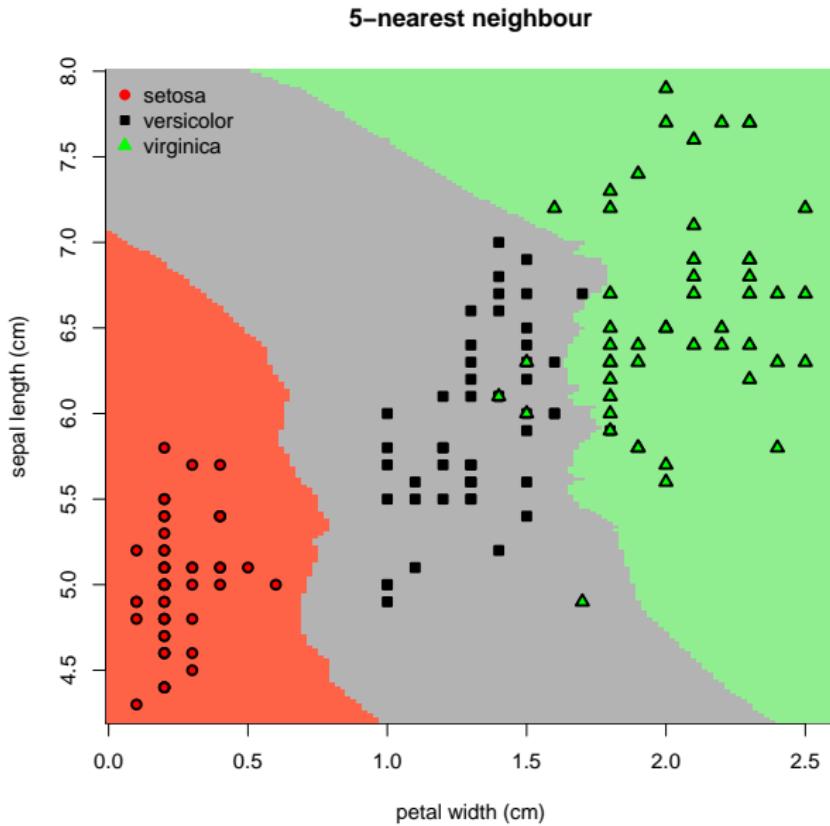
- ① Calculate distance between test point and every training data point
- ② Find the  $k$  training points closest to test point
- ③ Assign test point the majority vote of their class label

# $k$ -nearest neighbour

1-nearest neighbour

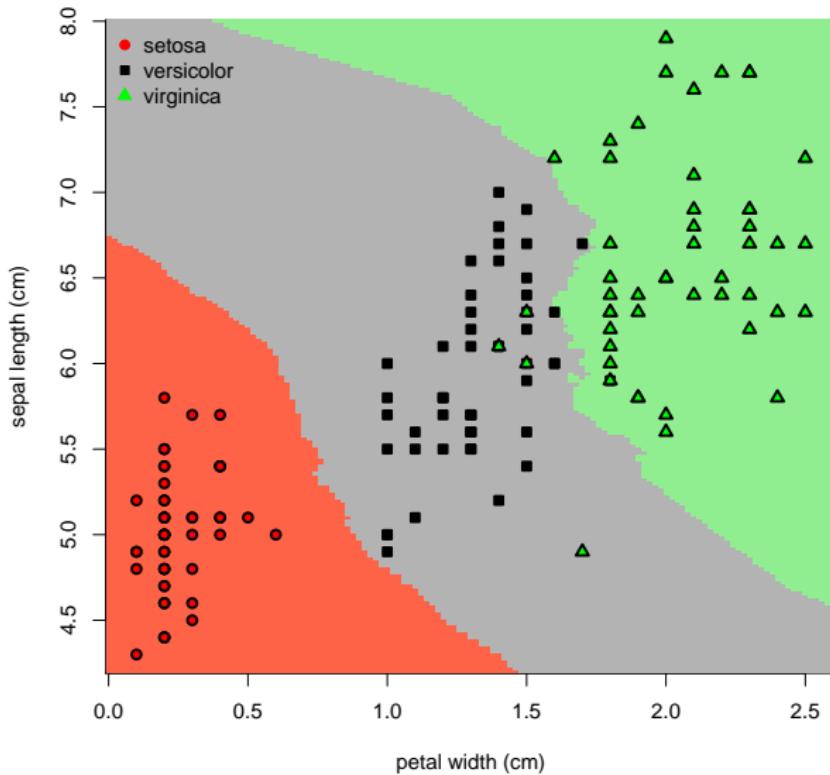


# $k$ -nearest neighbour



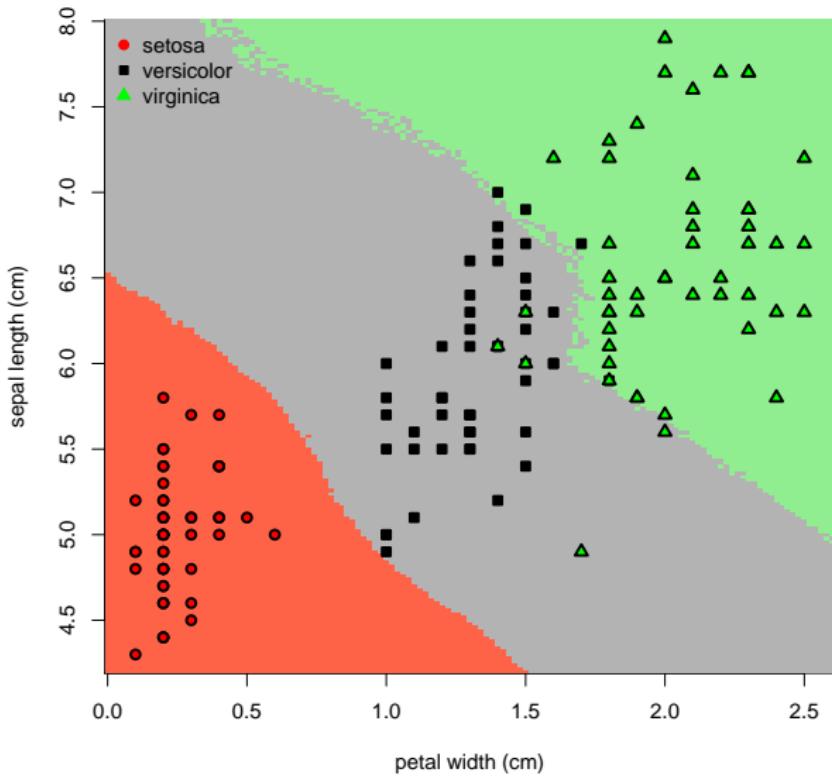
# $k$ -nearest neighbour

15-nearest neighbour



# $k$ -nearest neighbour

30-nearest neighbour



# $k$ -nearest neighbour

## Pros

- Simple and intuitive
- Works for multi-class problems
- Non-linear decision boundaries
- $k$  easily tuned by cross-validation

## Cons

- Can be computationally expensive, as for every test point, distance to every training data point needs to be computed
- Takes up a lot of storage as *all* training points need to be retained

# Decision trees

- ① Find the yes/no rule that best splits the data with respect to *one* of the features
- ② The best split is the one that produces the most homogeneous groups; found by maximising information gain/lowering entropy.
- ③ Repeat steps 1 to 2 until all data are correctly classified or some stopping rule reached.

# Decision trees

- ① Find the yes/no rule that best splits the data with respect to *one* of the features
- ② The best split is the one that produces the most homogeneous groups; found by maximising information gain/lowering entropy.
- ③ Repeat steps 1 to 2 until all data are correctly classified or some stopping rule reached.

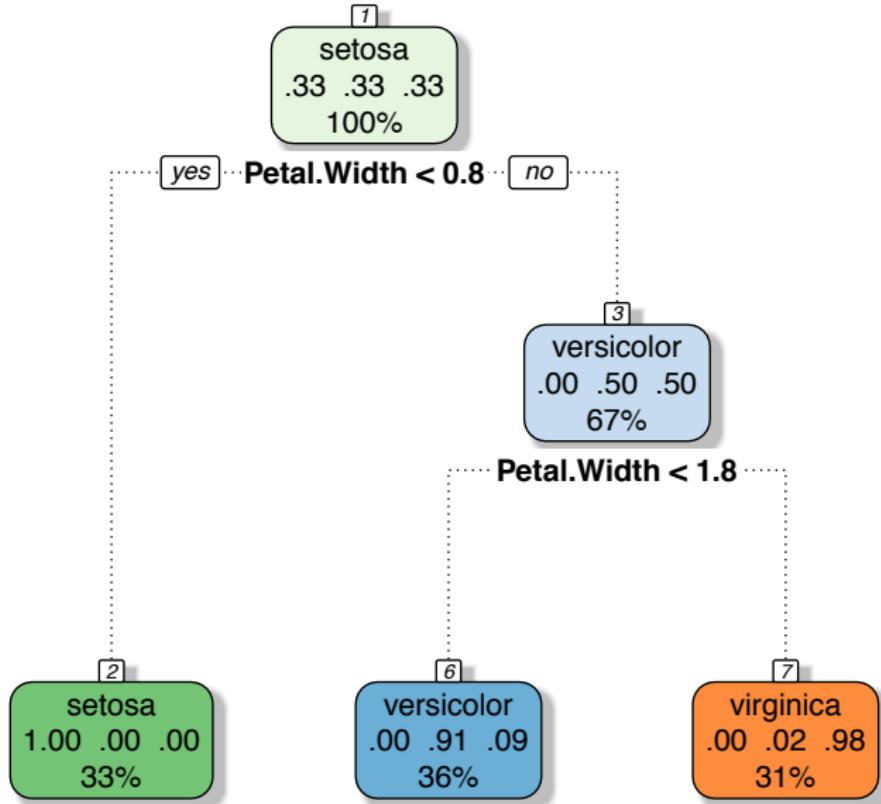
# Decision trees

- ① Find the yes/no rule that best splits the data with respect to *one* of the features
- ② The best split is the one that produces the most homogeneous groups; found by maximising information gain/lowering entropy.
- ③ Repeat steps 1 to 2 until all data are correctly classified or some stopping rule reached.

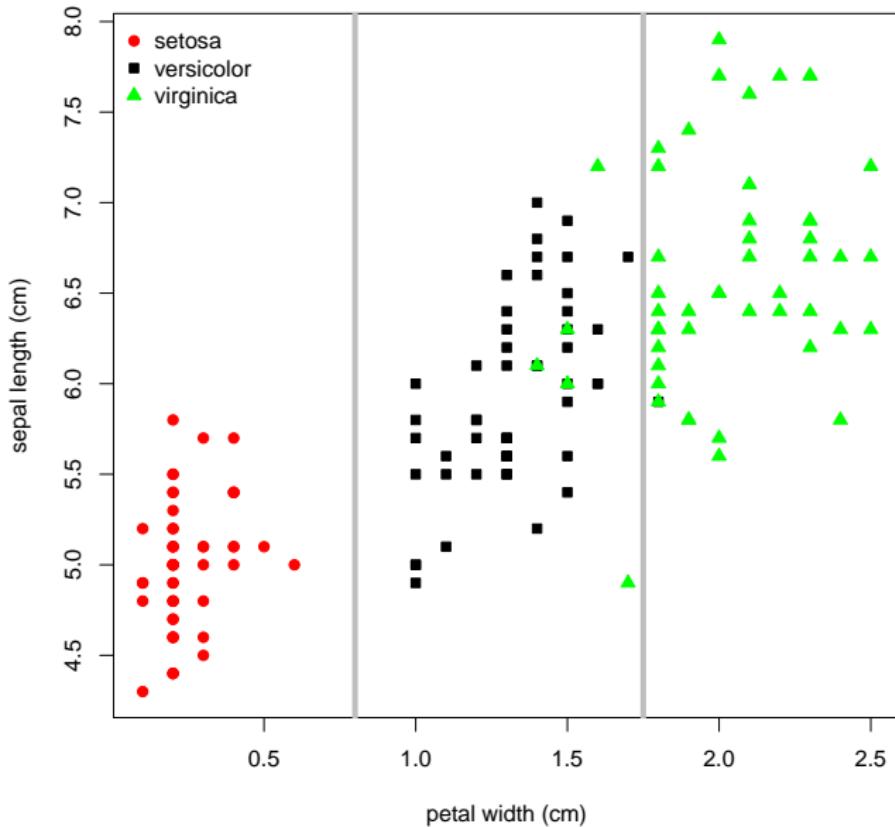
# Decision trees

- ① Find the yes/no rule that best splits the data with respect to *one* of the features
- ② The best split is the one that produces the most homogeneous groups; found by maximising information gain/lowering entropy.
- ③ Repeat steps 1 to 2 until all data are correctly classified or some stopping rule reached.

# Decision trees



# Decision boundaries



# Decision trees

## Pros

- Model is very easy to explain to non-experts and can be directly used to generate rules
- Computationally inexpensive to train, evaluate and store
- Handle both categorical and continuous data
- Robust to outliers

## Cons

- Can easily overfit the data
- Predictive accuracy can be poor
- Linear decision boundaries
- Small changes to training data may lead to a completely different tree

# Random forests

- Decision trees are intuitive but suffer from overfitting which significantly affect their predictive accuracy
- **Pruning**, to “trim” the tree back, help reduce this overfit
- **Ensemble** methods such as Random Forests are a better alternative
- **Rationale:** Instead of one tree, grow a *forest*, where every bushy tree (no pruning) is a bit different, then average predictions over all trees

# Random forests

- Decision trees are intuitive but suffer from overfitting which significantly affect their predictive accuracy
- Pruning, to “trim” the tree back, help reduce this overfit
- Ensemble methods such as Random Forests are a better alternative
- Rationale: Instead of one tree, grow a *forest*, where every bushy tree (no pruning) is a bit different, then average predictions over all trees

# Random forests

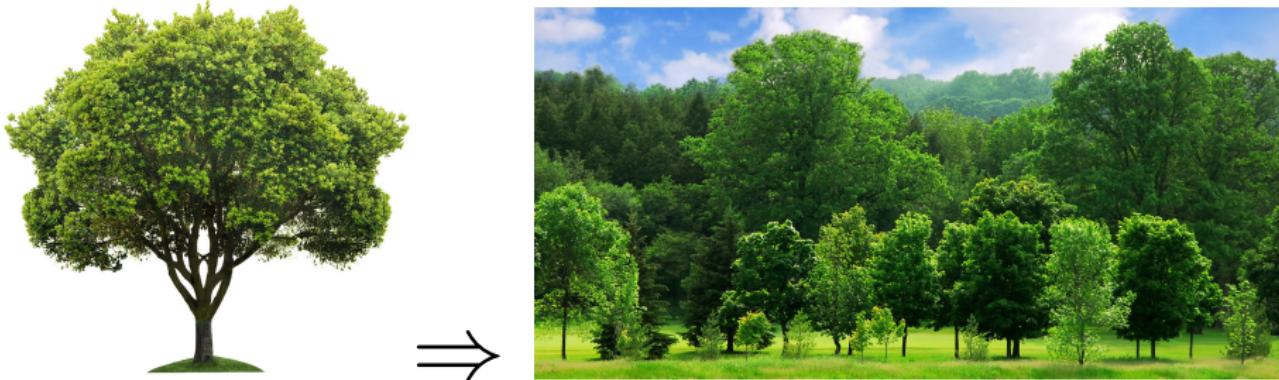
- Decision trees are intuitive but suffer from overfitting which significantly affect their predictive accuracy
- **Pruning**, to “trim” the tree back, help reduce this overfit
- Ensemble methods such as Random Forests are a better alternative
- **Rationale:** Instead of one tree, grow a *forest*, where every bushy tree (no pruning) is a bit different, then average predictions over all trees

# Random forests

- Decision trees are intuitive but suffer from overfitting which significantly affect their predictive accuracy
- **Pruning**, to “trim” the tree back, help reduce this overfit
- **Ensemble** methods such as Random Forests are a better alternative
  - **Rationale:** Instead of one tree, grow a *forest*, where every bushy tree (no pruning) is a bit different, then average predictions over all trees

# Random forests

- Decision trees are intuitive but suffer from overfitting which significantly affect their predictive accuracy
- **Pruning**, to “trim” the tree back, help reduce this overfit
- **Ensemble** methods such as Random Forests are a better alternative
- **Rationale:** Instead of one tree, grow a *forest*, where every bushy tree (no pruning) is a bit different, then average predictions over all trees



# Random forests

① Grow  $T$  decorrelated trees (no pruning)

② Forest randomness is induced by:

③ Average predictions from all  $T$  trees.

# Random forests

① Grow  $T$  decorrelated trees (no pruning)

② Forest randomness is induced by:

③ Average predictions from all  $T$  trees.

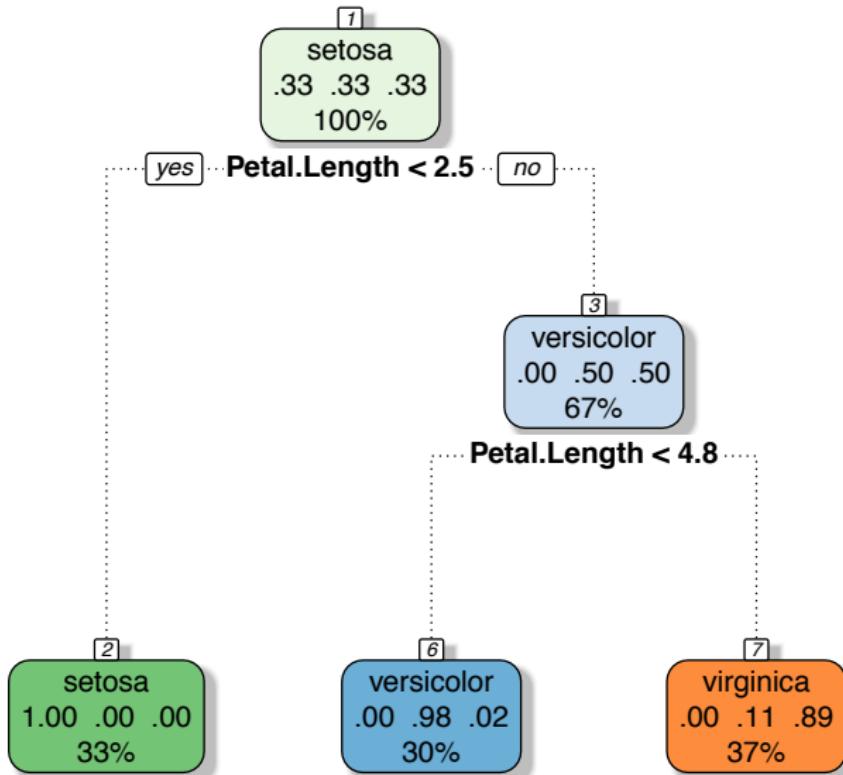
# Random forests

- ① Grow  $T$  decorrelated trees (no pruning)
- ② Forest randomness is induced by:
  - Bagging (**Bootstrap AGGregatING**), each tree is trained on a subset of the data randomly sampled with replacement
  - Considering only a subset of predictors as candidates for each split
- ③ Average predictions from all  $T$  trees.

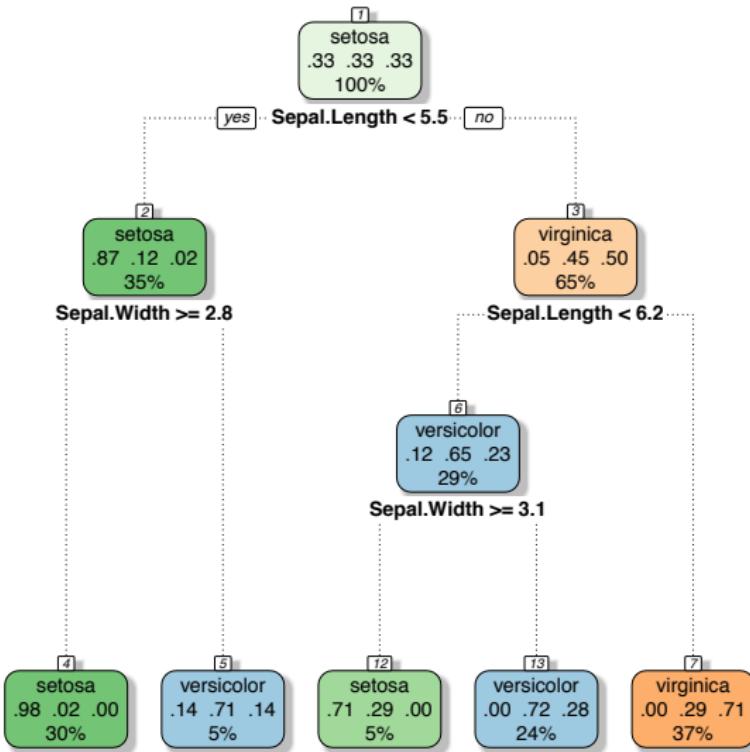
# Random forests

- ① Grow  $T$  decorrelated trees (no pruning)
- ② Forest randomness is induced by:
  - Bagging (**Bootstrap AGGregatING**), each tree is trained on a subset of the data randomly sampled with replacement
  - Considering only a subset of predictors as candidates for each split
- ③ Average predictions from all  $T$  trees.

# De-correlated trees

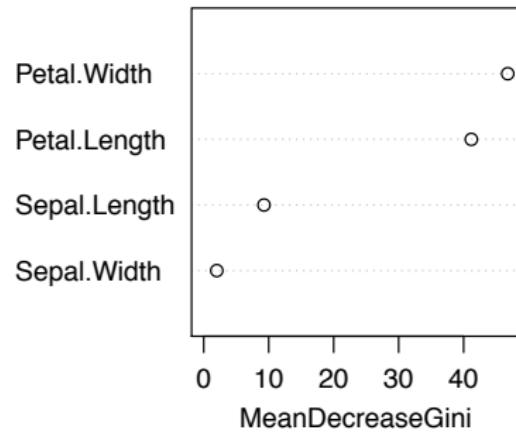
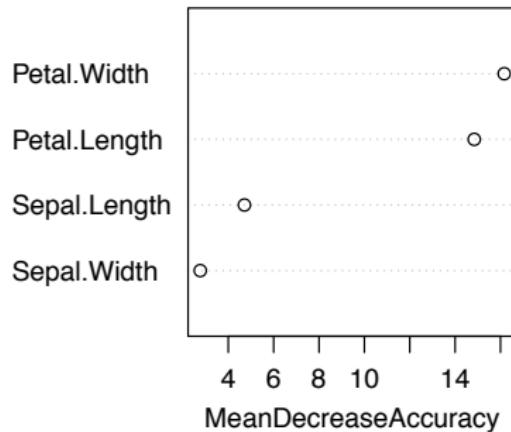


# De-correlated trees



# Variable importance

- Cannot visualise decision boundaries (loss of interpretability)
- However, variable importance helps us perform feature selection



# Random forests

## Pros

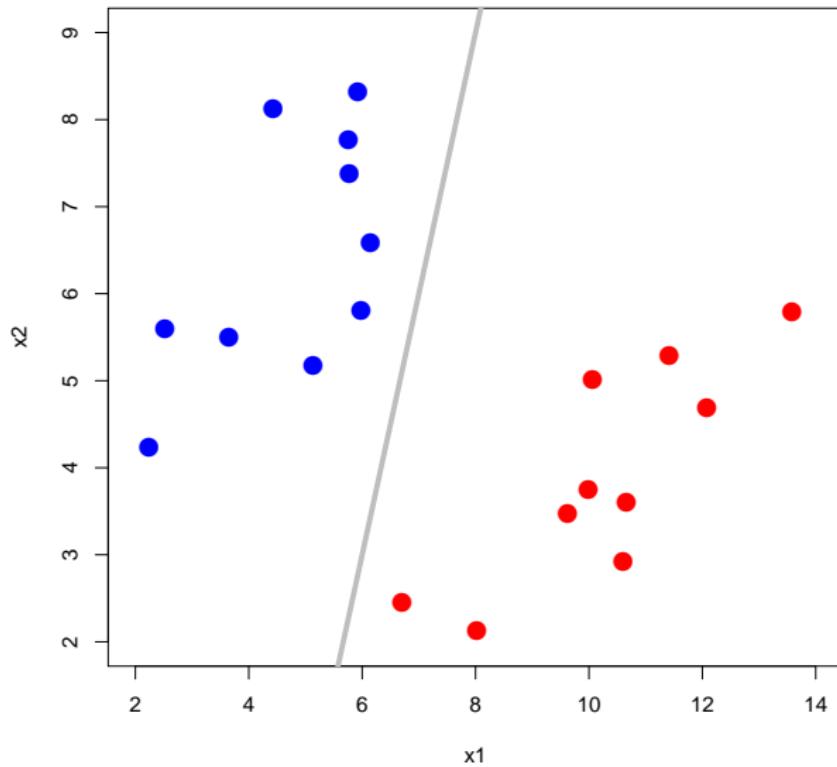
- State-of-the-art predictive accuracy
- Can handle thousands of both categorical and continuous predictors without variable deletion
- Robust to outliers
- Estimates the importance of every predictor
- Out-of-bag error (unbiased estimate of test error for every tree built)
- Can cope with unbalanced datasets by setting class weights
- Trivially parallelisable

## Cons

- Harder to interpret than plain decision trees

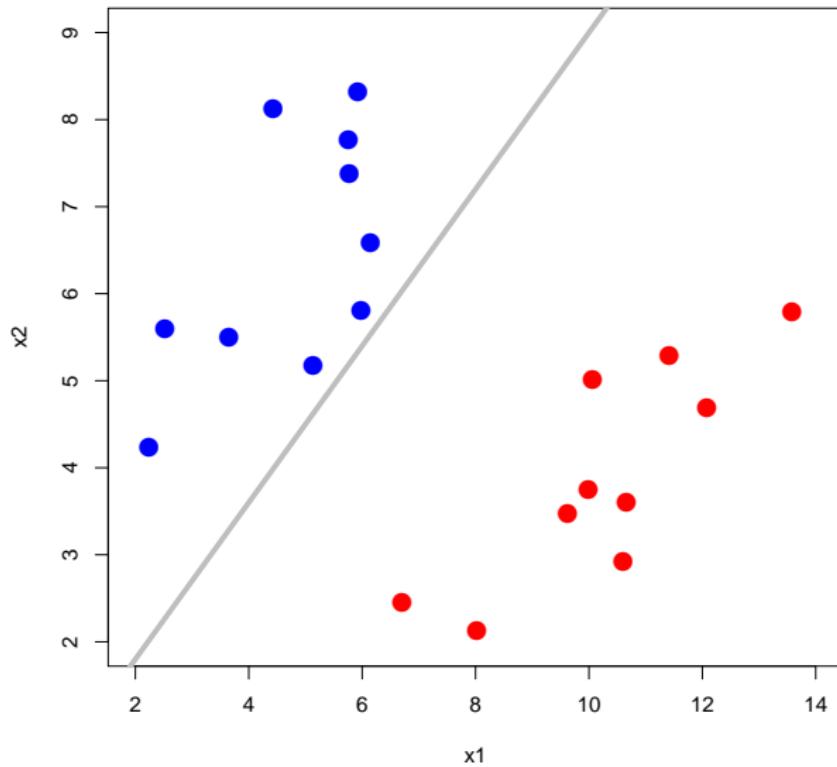
# Support vector machines (SVMs)

Which is the best separating line?



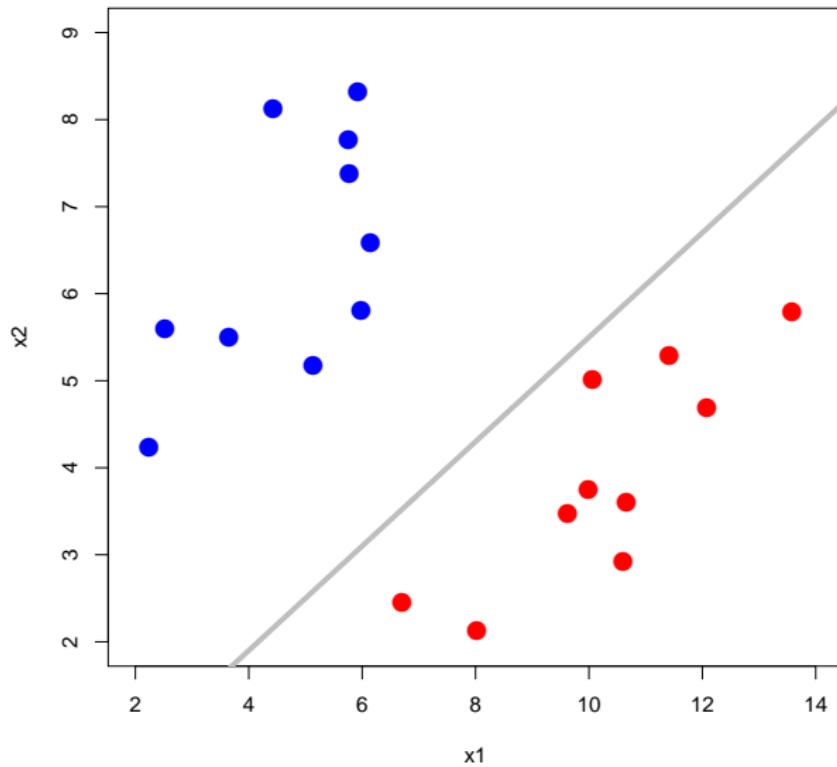
# Support vector machines (SVMs)

Which is the best separating line?



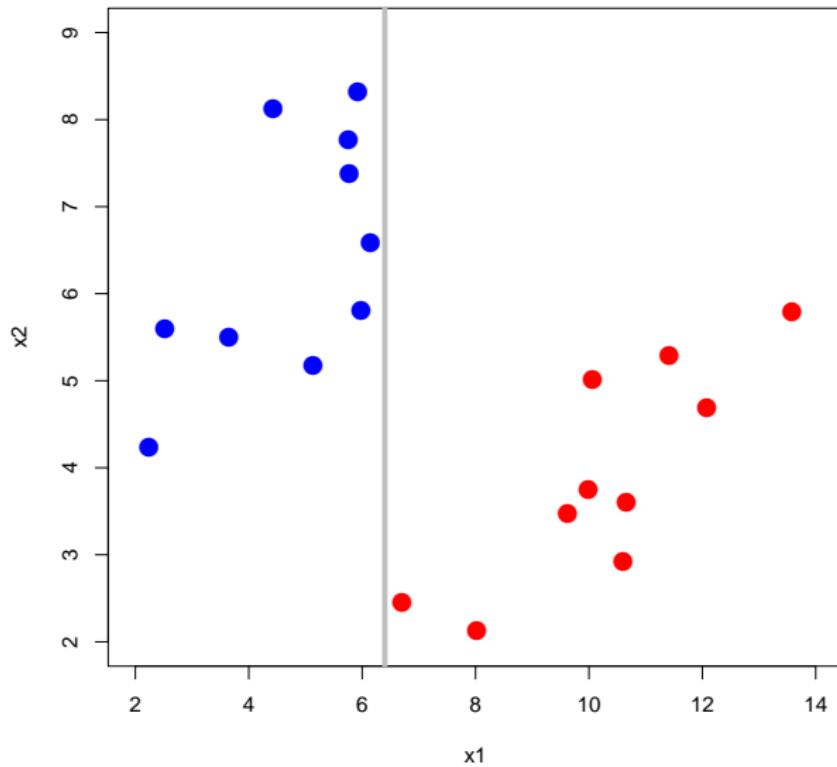
# Support vector machines (SVMs)

Which is the best separating line?



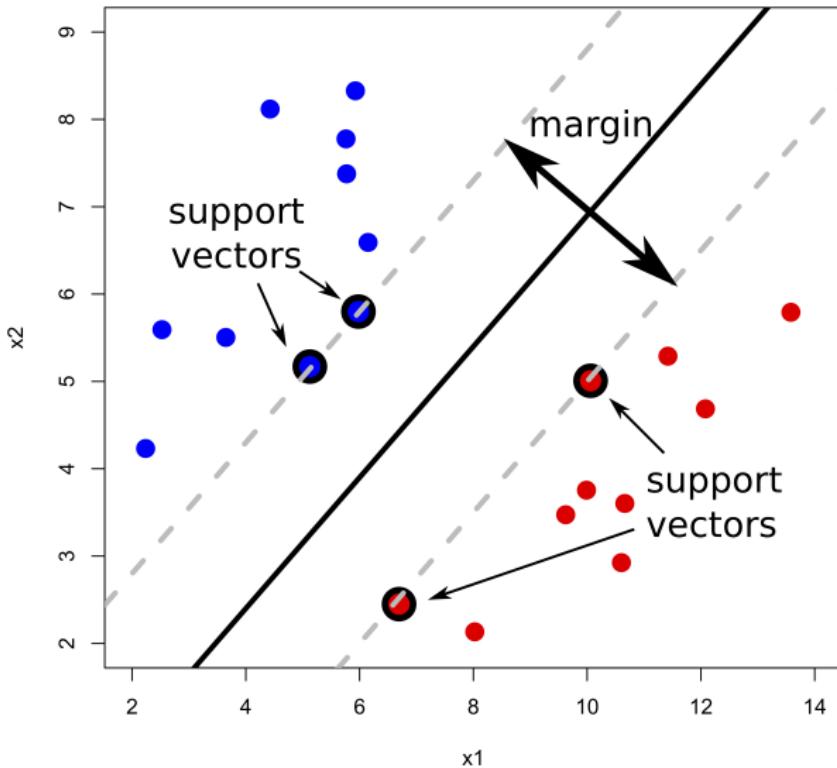
# Support vector machines (SVMs)

Which is the best separating line?



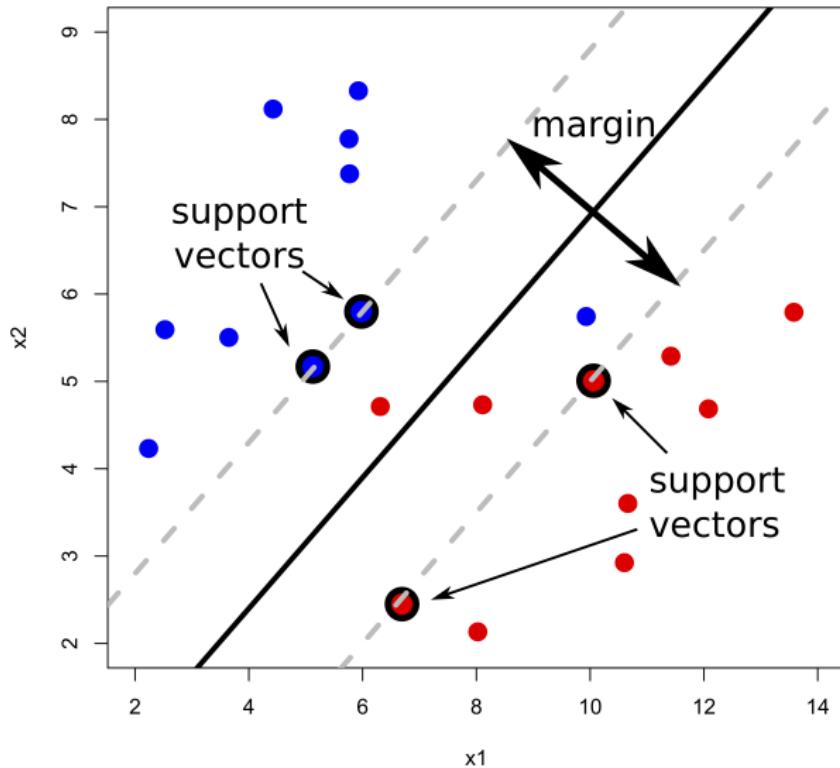
# Maximal margin classifier

**Rationale:** Maximise the *margin*



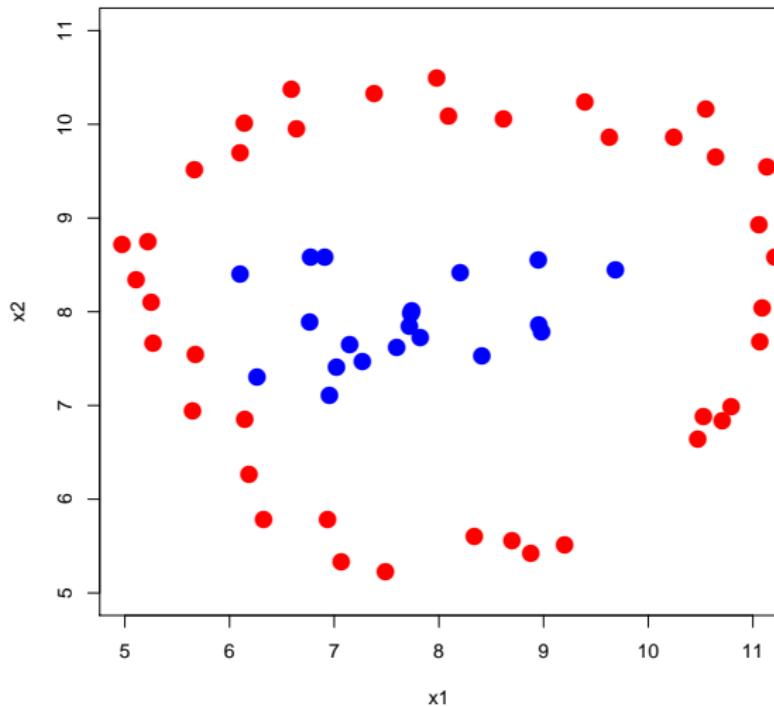
# Support vector classifiers

**Rationale:** Use a *soft* margin



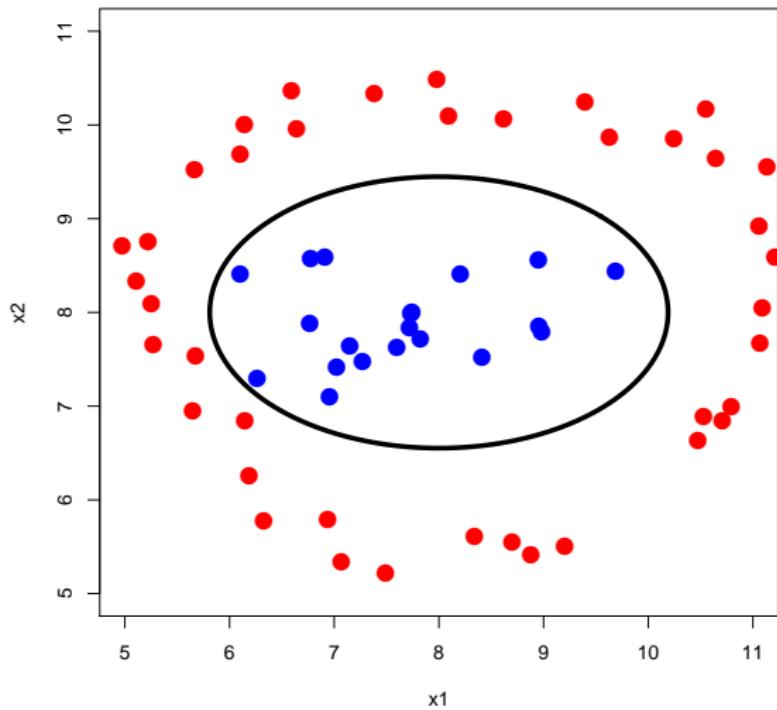
# Support vector machines

Real-life data is complex and often we cannot find a separating hyperplane



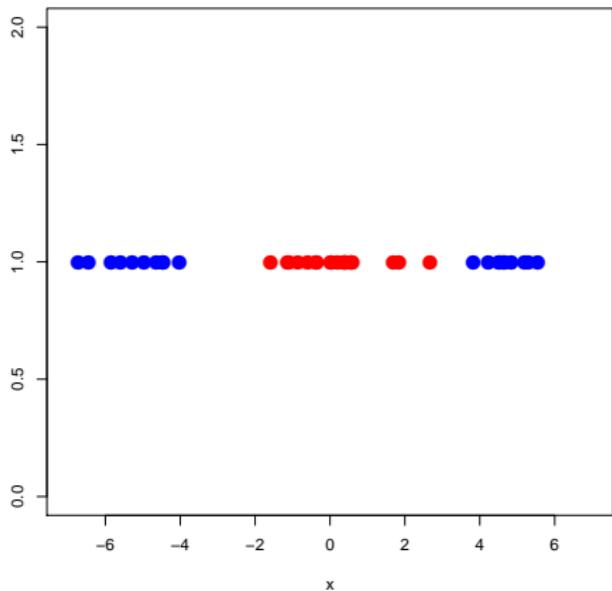
# Support vector machines

**Rationale:** Map data to a higher dimensional space where classes are linearly separable



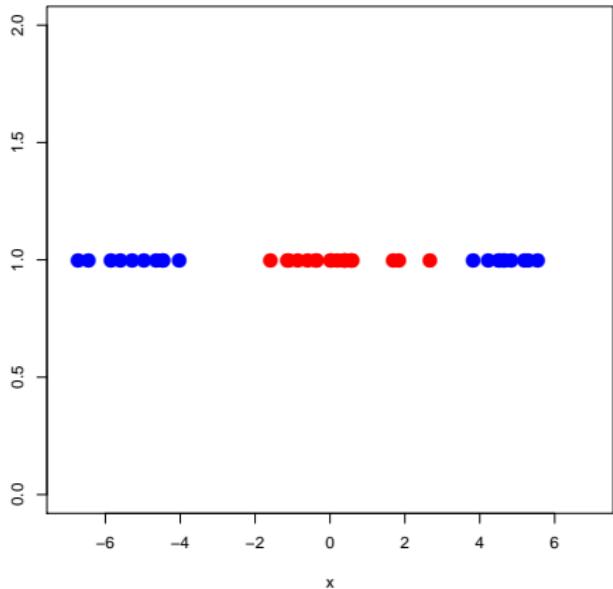
# Support vector machines: 1D to 2D

1D (original) data is not linearly separable

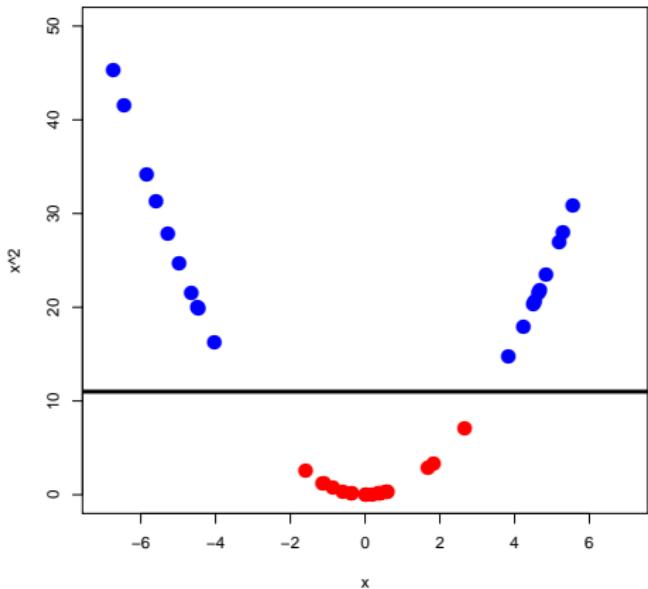


# Support vector machines: 1D to 2D

1D (original) data is not linearly separable



2D (transformed) data is now linearly separable



# Support vector machines - the kernel trick

- So our solution is to blow up the dimensions?
- But what about the “curse of dimensionality”?
- Very computationally expensive to work in high dimensions

# Support vector machines - the kernel trick

- So our solution is to blow up the dimensions?
- But what about the “curse of dimensionality”?
- Very computationally expensive to work in high dimensions

- **Kernel trick** to the rescue!
- Work in an *implicit* feature space
- Data is never explicitly computed in higher dimensions
- Think about kernels as generalised distance measures

p.s Kernel methods are mathematically intricate and beyond the scope of this introductory workshop

# Support vector machines

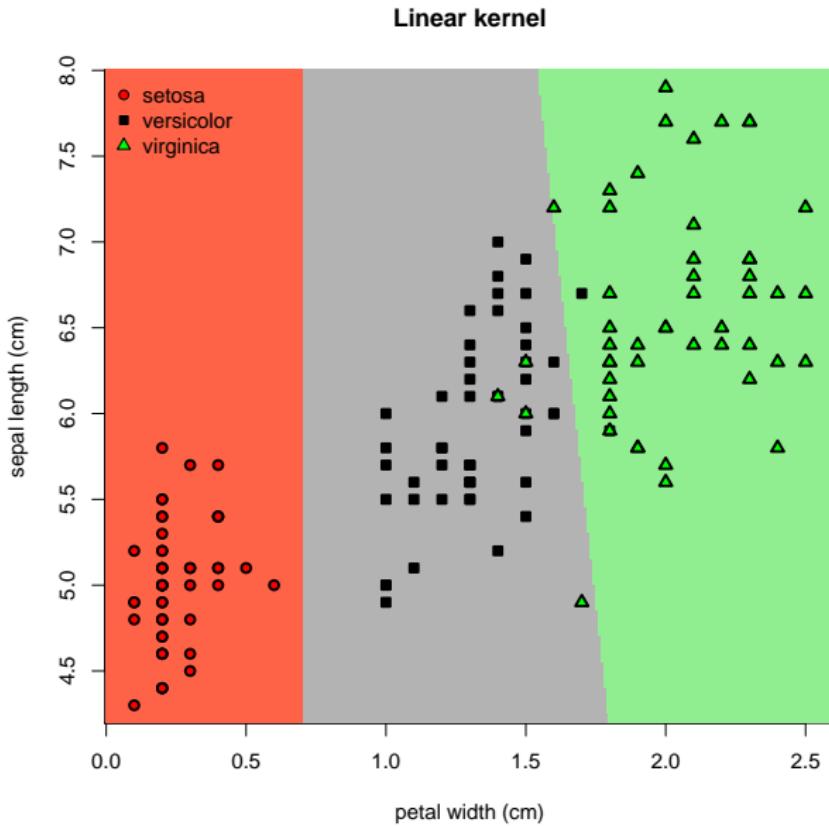
- ① Choose a kernel
- ② Run optimiser to find the maximum margin separating hyperplane

# Support vector machines

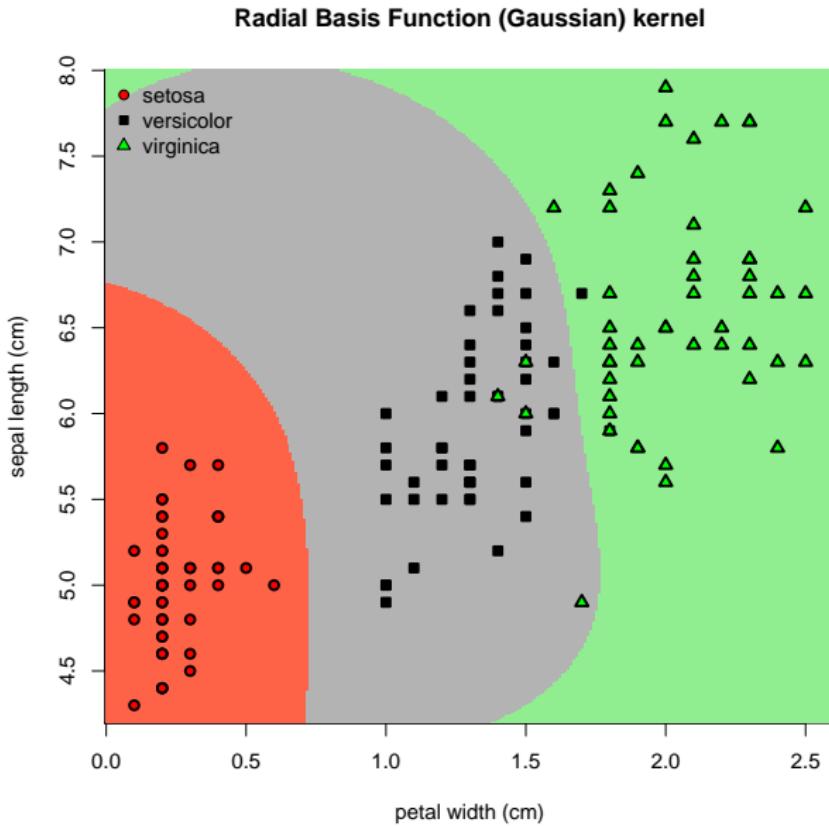
- ① Choose a kernel
- ② Run optimiser to find the maximum margin separating hyperplane

SVMs are inherently binary classifiers. The most common ways to deal with multi-class problems is by building several **one-versus-all** or **one-versus-one** classifiers.

# Support vector machines



# Support vector machines



# Support vector machines

## Pros

- State-of-the-art predictive accuracy
- Low storage requirements (only support vectors to store)
- A vast array of kernels are available that are flexible enough to cater for any type of data
- Global optimum guaranteed

## Cons

- Model is hard to interpret
- Feature space cannot be visualised